

TensorFlow For Practice

By: Mohamed Aziz Tousli

Basics

► #Necessary libraries

```
import tensorflow as tf
```

```
from tensorflow import keras
```

► #Import an already existing dataset

```
data_set_name = keras.datasets.data_set_name
```

```
(train_images, train_labels), (test_images, test_labels) = data_set_name.load_data()
```

```
train_images = train_images / 255.0 #Normalization
```

► #Basic code

```
model = /* model definition here */
```

```
model.compile(optimizer='optimizer_name', loss='loss_name', metrics=['metrics_name'])
```

```
model.fit(x, y, epochs=number_of_epochs, callbacks=[callbacks])
```

```
model.evaluate(x_test, y_test)
```

```
model.predict([10])
```

```
model.summary() #Summary of model architecture
```

How to “learn”?

1- Make a guess.

2- The LOSS function measures the guessed answers against the known correct answers

3- The OPTIMIZER function makes another guess. Based on how the loss function went, it will try to minimize the loss.

4- It will repeat this for the number of EPOCHS.

Callback to stop training on condition

```
class myCallback(tf.keras.callbacks.Callback):  
    def on_epoch_end(self, epoch, logs={}):  
        if(logs.get('loss')<0.4):  
            print("\nLoss is low so cancelling training!")  
            self.model.stop_training = True
```

callbacks = myCallback()

→ model.fit(x, y, epochs=number_of_epochs, callbacks=[callbacks])

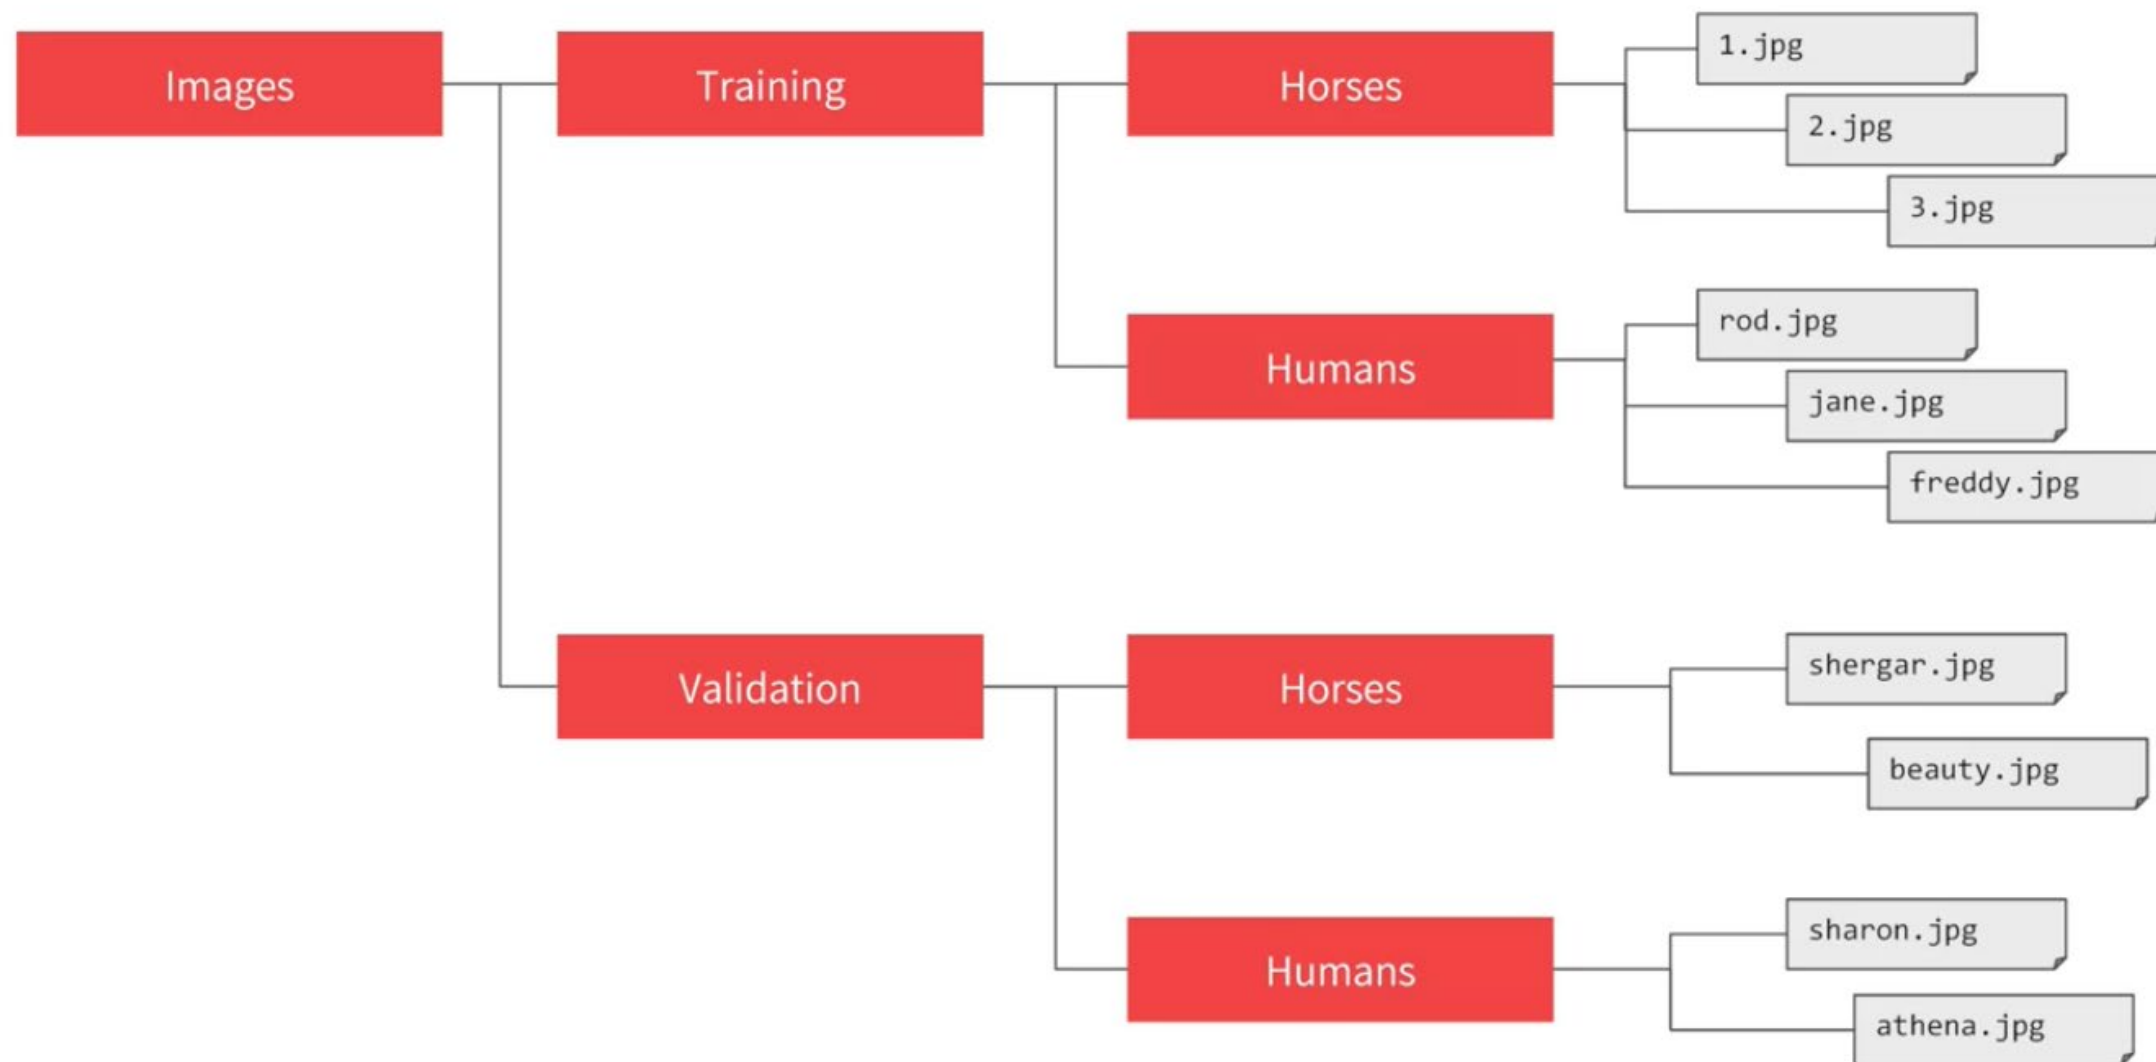
Model for Image Classification

```
model = keras.Sequential([ #Define a sequence of layers
    keras.layers.Flatten(input_shape=(x_pixel_size, y_pixel_size, 1|3)), #nD → 1D
    keras.layers.Dense(number_of_neurons, activation=tf.nn.relu), #First layer of neurones
    → If multi classification: keras.layers.Dense(number_of_classes, activation='softmax')
    → If binary classification: keras.layers.Dense(1, activation='sigmoid')
])
```

Model for Image Classification + Convolution

```
training_images = training_images.reshape(number_of_images, x_image_size, y_image_size, 1)
model = keras.Sequential([
    keras.layers.Conv2D(number_of_filters, (x_filter_size, y_filter_size), activation='relu',
                        input_shape=(x_image_size, y_image_size, 1 | 3)),
    keras.layers.MaxPooling2D((x_pooling_size, y_pooling_size),
                              kernel_size=(x_pooling_size, y_pooling_size)),
    keras.layers.Conv2D(number_of_filters, (x_filter_size, y_filter_size), activation='relu'),
    keras.layers.MaxPooling2D((x_pooling_size, y_pooling_size),
                              kernel_size=(x_pooling_size, y_pooling_size)),
    keras.layers.Flatten(),
    keras.layers.Dense(number_of_neurons, activation='relu'),
    keras.layers.Dense(number_of_classes, activation='softmax')])
```

ImageDataGenerator



ImageDataGenerator

- ▶ Goal: Read images from subdirectories, and automatically label them

```
from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
train_datagen = ImageDataGenerator(rescale=1./255)
```

```
train_generator = train_datagen.flow_from_directory(
```

```
    train_dir,
```

```
    target_size=(x_image_size, y_image_size),
```

```
    batch_size=size_of_batch,
```

```
    class_mode='binary' | 'categorical')
```

- ▶ PS: Same thing has to be done for validation_generator

- ▶ fit → fit_generator | evaluate → evaluate_generator | predict → predict_generator

```
history = model.fit_generator(
```

```
    train_generator, steps_per_epoch=8, epochs=15, verbose=2,
```

```
    validation_data=validation_generator, validation_steps=8)
```

Prediction with upload button

```
import numpy as np
from google.colab import files
from keras.preprocessing import image

uploaded = files.upload()

for fn in uploaded.keys():

    # predicting images
    path = '/content/' + fn
    img = image.load_img(path, target_size=(300, 300))
    x = image.img_to_array(img)
    x = np.expand_dims(x, axis=0)

    images = np.vstack([x])
    classes = model.predict(images, batch_size=10)
    print(classes[0])
    if classes[0]>0.5:
        print(fn + " is a human")
    else:
        print(fn + " is a horse")
```


Data Augmentation with ImageDataGenerator

```
# Updated to do image augmentation
train_datagen = ImageDataGenerator(
    rotation_range=40,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest')
```

These are just a few of the options available (for more, see the Keras documentation. Let's quickly go over what we just wrote:

- `rotation_range` is a value in degrees (0–180), a range within which to randomly rotate pictures.
- `width_shift` and `height_shift` are ranges (as a fraction of total width or height) within which to randomly translate pictures vertically or horizontally.
- `shear_range` is for randomly applying shearing transformations.
- `zoom_range` is for randomly zooming inside pictures.
- `horizontal_flip` is for randomly flipping half of the images horizontally. This is relevant when there are no assumptions of horizontal assymetry (e.g. real-world pictures).
- `fill_mode` is the strategy used for filling in newly created pixels, which can appear after a rotation or a width/height shift.

Transfer Learning

- ▶ `from tensorflow.keras import Model`
- ▶ `from tensorflow.keras.applications.inception_v3 import InceptionV3`
- ▶ `local_weights_file = 'h_5_file'`
- ▶ `pre_trained_model = InceptionV3(input_shape=(x,y,1|3), include_top=False, weights=None)`
- ▶ `pre_trained_model.load_weights(local_weights_file)`
- ▶ `for layer in pre_trained_model.layers:`
 - `layer.trainable = False`
- ▶ `last_layer = pre_trained_model.get_layer('last_layer_name')` #Get last layer of Inception
- ▶ `last_output = get_layer.output`
- ▶ `x = layers.Flatten()(last_output)` #Work from last layer of Inception (Transfer Learning)
- ▶ `x = layers.Dense(1024, activation='relu')(x)`
- ▶ `x = layers.Dense(1, activation='sigmoid')(x)`
- ▶ `x = layers.Dropout(0.2)(x)`
- ▶ `model = Model(pre_trained_model.input, x)`

NLP

- ▶ `from tensorflow.keras.preprocessing.text import Tokenizer` #Tokenizer
- ▶ `from tensorflow.keras.preprocessing.sequence import pad_sequences` #Padding to the longest sentence
- ▶ `sentences = ['sentence1', 'sentence2']`
- ▶ `tokenizer = Tokenizer(num_words = vocab_size, oov_token="<OOV>")`
- ▶ `tokenizer.fit_on_texts(sentences)`
- ▶ `word_index = tokenizer.word_index` # `word_index = {'a': 1, 'dog': 2, ...}`
- ▶ `sequences = tokenizer.texts_to_sequences(sentences)` # `sequences = [[4,2,3,1], ...]`
- ▶ `padded = pad_sequences(sequences)` # `padded = [[0,0,0,4,2,3,1], ...]`
 - ▶ Other arguments: `padding='post'`, `maxlen=max_length`, `truncating=trunc_type`
- ▶ `tokenized_string = tokenizer.encode('sample_string')`
- ▶ `original_string = tokenizer.decode(tokenized_string)`

NLP with a Dataset

- ▶ `import tensorflow_datasets as tfds`
 - ▶ `db_name, info = tfds.load("db_name", with_info=True, as_supervised=True)`
 - ▶ `train_data, test_data = db_name['train'], db_name['test']`
 - ▶ `tokenizer = info.features['text'].encoder`
- ▶ `training_sentences, training_labels, testing_sentences, testing_labels = [], [], [], []`
- ▶ `for s,l in train_data:`
 - `training_sentences.append(str(s.tonumpy()))`
 - `training_labels.append(l.numpy())`
- ▶ `for s,l in test_data:`
 - `testing_sentences.append(str(s.tonumpy()))`
 - `testing_labels.append(l.numpy())`
- ▶ `model = tf.keras.Sequential([`
 - `tf.keras.layers.Embedding(vocab_size, embedding_dim, input_length=max_length),`
 - `tf.keras.layers.GlobalAveragePooling1D(),` `#tf.keras.layers.Flatten()` ← We can use this instead
 - `tf.keras.layers.Dense(6, activation='relu'),`
 - `tf.keras.layers.Dense(1, activation='sigmoid'))]`
- ▶ `model.fit(training_padded, training_labels, epochs=num_epochs, validation_data=(testing_padded, testing_labels))`

NLP with LSTM | GRU

- ▶ `model = tf.keras.Sequential([`
 `tf.keras.layers.Embedding(tokenizer.vocab_size, 64),`
 `tf.keras.layers.Bidirectional(tf.keras.layers.LSTM|GRU(64, return_sequences=True)),`
 `tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(64)),` #If we add this layer → `return_sequences=True`
 `tf.keras.layers.Dense(64, activation='relu'),`
 `tf.keras.layers.Dense(1, activation='sigmoid'))]`
- ▶ PS:
 - ▶ We can use `Bidirectional()` instead of `Flatten()` & `GlobalAveragePooling1D()`
 - ▶ We can use `Conv1D()` & `GlobalMaxPooling1D()` instead of `Bidirectional()`

Predict Next Word {1} - Preprocessing

```
input_sequences = []
for line in corpus:
    token_list = tokenizer.texts_to_sequences([line])[0]
    for i in range(1, len(token_list)):
        n_gram_sequence = token_list[:i+1]
        input_sequences.append(n_gram_sequence)
max_sequence_len = max([len(x) for x in input_sequences])
```

Padded Input Sequences:

Input (X) → [0 0 0 0 0 0 0 0 0 4 2] 66 → Label (Y)

[0 0 0 0 0 0 0 0 0 4 2]
[0 0 0 0 0 0 0 0 4 2 66 8]
[0 0 0 0 0 0 0 4 2 66 8 67]
[0 0 0 0 0 0 4 2 66 8 67 68]
[0 0 0 0 0 4 2 66 8 67 68 69]
[0 0 0 0 4 2 66 8 67 68 69 70]

Predict Next Word {2} - Preprocessing

```
xs = input_sequences[:, :-1]
labels = input_sequences[:, -1]

ys = tf.keras.utils.to_categorical(labels, num_classes=total_words)
```

Sentence: [0 0 0 0 4 2 66 8 67 68 69 70]

```
x:[0 0 0 0 4 2 66 8 67 68 69]
```

Label: [70]

Predict Next Word {2} - Training

```
model = Sequential()  
model.add(Embedding(total_words, 100, input_length=max_sequence_len-1))  
model.add(Bidirectional(LSTM(150)))  
model.add(Dense(total_words, activation='softmax'))  
adam = Adam(lr=0.01)  
model.compile(loss='categorical_crossentropy', optimizer=adam, metrics=['accuracy'])  
history = model.fit(xs, ys, epochs=100, verbose=1)
```


Predict Next Word {2} - Predicting

```
seed_text = "Laurence went to dublin"
next_words = 10

for _ in range(next_words):
    token_list = tokenizer.texts_to_sequences([seed_text])[0]
    token_list = pad_sequences([token_list], maxlen=max_sequence_len - 1, padding='pre')
    predicted = model.predict_classes(token_list, verbose=0)
    output_word = ""
    for word, index in tokenizer.word_index.items():
        if index == predicted:
            output_word = word
            Break
    seed_text += " " + output_word
print(seed_text)
```

Time Series

► Notebooks:

- Introduction to time series:

https://colab.research.google.com/github/lmoroney/dlaicourse/blob/master/TensorFlow%20In%20Practice/Course%204%20-%20S%2BP/S%2BP_Week_1_Lesson_2.ipynb

- Forecasting:

<https://colab.research.google.com/github/lmoroney/dlaicourse/blob/master/TensorFlow%20In%20Practice/Course%204%20-%20S%2BP/S%2BP%20Week%201%20-%20Lesson%203%20-%20Notebook.ipynb>

- Preparing features and labels:

<https://colab.research.google.com/github/lmoroney/dlaicourse/blob/master/TensorFlow%20In%20Practice/Course%204%20-%20S%2BP/S%2BP%20Week%202%20Lesson%201.ipynb>

- Single layer neural network:

<https://colab.research.google.com/github/lmoroney/dlaicourse/blob/master/TensorFlow%20In%20Practice/Course%204%20-%20S%2BP/S%2BP%20Week%202%20Lesson%202.ipynb>

- Deep neural network:

<https://colab.research.google.com/github/lmoroney/dlaicourse/blob/master/TensorFlow%20In%20Practice/Course%204%20-%20S%2BP/S%2BP%20Week%202%20Lesson%203.ipynb>

Time Series

► Notebooks:

► RNN:

<https://colab.research.google.com/github/lmoroney/dlaicourse/blob/master/TensorFlow%20In%20Practice/Course%204%20-%20S%2BP/S%2BP%20Week%203%20Lesson%202%20-%20RNN.ipynb>

► LSTM1:

<https://colab.research.google.com/github/lmoroney/dlaicourse/blob/master/TensorFlow%20In%20Practice/Course%204%20-%20S%2BP/S%2BP%20Week%203%20Lesson%204%20-%20LSTM.ipynb>

► LSTM2:

<https://colab.research.google.com/github/lmoroney/dlaicourse/blob/master/TensorFlow%20In%20Practice/Course%204%20-%20S%2BP/S%2BP%20Week%204%20Lesson%201.ipynb>

► Sunspots1:

<https://colab.research.google.com/github/lmoroney/dlaicourse/blob/master/TensorFlow%20In%20Practice/Course%204%20-%20S%2BP/S%2BP%20Week%204%20Lesson%205.ipynb>

► Sunspots2:

<https://colab.research.google.com/github/lmoroney/dlaicourse/blob/master/TensorFlow%20In%20Practice/Course%204%20-%20S%2BP/S%2BP%20Week%204%20Lesson%203.ipynb>